

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
```

```
In [ ]:
```

```
In [2]: import pandas as pd

# Read the CSV file into a DataFrame
data = pd.read_csv('your_file.csv')

from sklearn.preprocessing import LabelEncoder

# Initialize the LabelEncoder
label_encoder = LabelEncoder()

# Fit and transform the 'sentiment' column
data['sentiment_encoded'] = label_encoder.fit_transform(data['Sentiment'])

# Display the DataFrame with the encoded sentiment column
data
```

Out[2]:

	Product Name	Brand Name	Price	Rating	Reviews	Review Votes	Sentiment	Tokenized	Without_Stop
0	"clear clean esn" sprint epic 4g galaxy sph-d7...	samsung	199.99	5	i feel so lucky to have found this used (phone...	1.0	positive	['i', 'feel', 'so', 'lucky', 'to', 'have', 'fo...]	['feel', 'found', 'u:']
1	"clear clean esn" sprint epic 4g galaxy sph-d7...	samsung	199.99	4	nice phone, nice up grade from my pantach revu...	0.0	positive	['nice', 'phone', ',', 'nice', 'up', 'grade', ...]	['nice', 'pho', 'nice', ',']
2	"clear clean esn" sprint epic 4g galaxy sph-d7...	samsung	199.99	5	very pleased	0.0	positive	['very', 'pleased']	['pl']
3	"clear clean esn" sprint epic 4g galaxy sph-d7...	samsung	199.99	4	it works good but it goes slow sometimes but i...	0.0	positive	['it', 'works', 'good', 'but', 'it', 'goes', '...]	['works', 'goes', 'somet']
4	"clear clean esn" sprint epic 4g galaxy sph-d7...	samsung	199.99	4	great phone to replace my lost phone. the only...	0.0	positive	['great', 'phone', 'to', 'replace', 'my', 'los...]	['great', 'r', 'replace', 'pl']
...	...	...	...	...	...	...	...	...	
349465	samsung convoy u640 phone for verizon wireless...	samsung	79.95	5	great phone. large keys, best flip phone i hav...	0.0	positive	['great', 'phone', ',', 'large', 'keys', ',', '...', ...]	['great', 'pho', 'large', 'key']
349466	samsung convoy u640 phone for verizon wireless...	samsung	79.95	5	pros...works great, very durable, easy to navi...	0.0	positive	['pros', '...', 'works', 'great', ',', 'very', ...]	['pros', '...', 'c', 'great', ',', 'c']
349467	samsung convoy	samsung	79.95	5	just as described	0.0	positive	['just', 'as', 'described',]	['described', 'p']

	Product Name	Brand Name	Price	Rating	Reviews	Review Votes	Sentiment	Tokenized	Without_Stop
	for verizon wireless...				perfect for the price			'perfect', 'for', ...	
349468	samsung convoy u640 phone for verizon wireless...	samsung	79.95	1	would not work	0.0	negative	['would', 'not', 'work']	['would',
349469	samsung convoy u640 phone for verizon wireless...	samsung	79.95	3	speaker phone doesn't work, but phone works good	0.0	neutral	['speaker', 'phone', 'does', 'n't', 'work', '...',	['speaker', 'l "n't", 'w



```
In [3]: data.head()
```

Out[3]:

	Product Name	Brand Name	Price	Rating	Reviews	Review Votes	Sentiment	Tokenized	Without_Stopwords
0	"clear clean esn" sprint epic 4g galaxy sph-d7...	samsung	199.99	5	i feel so lucky to have found this used (phone...	1.0	positive	['i', 'feel', 'so', 'lucky', 'to', 'have', 'found', 'this', 'used', 'phone...']	['feel', 'lucky', 'found', 'used', 'phone...']
1	"clear clean esn" sprint epic 4g galaxy sph-d7...	samsung	199.99	4	nice phone, nice up grade from my pantach revu...	0.0	positive	['nice', 'phone', 'up', 'grade', 'from', 'my', 'pantach', 'revu...']	['nice', 'phone', 'up', 'grade', 'panta...']
2	"clear clean esn" sprint epic 4g galaxy sph-d7...	samsung	199.99	5	very pleased	0.0	positive	['very', 'pleased']	['pleased']
3	"clear clean esn" sprint epic 4g galaxy sph-d7...	samsung	199.99	4	it works good but it goes slow sometimes but i...	0.0	positive	['it', 'works', 'good', 'but', 'it', 'goes', 'slow', 'sometimes', 'but', 'i...']	['works', 'good', 'goes', 'slow', 'sometimes', 'but', 'i...']
4	"clear clean esn" sprint epic 4g galaxy sph-d7...	samsung	199.99	4	great phone to replace my lost phone. the only...	0.0	positive	['great', 'phone', 'to', 'replace', 'my', 'lost', 'phone', 'the', 'only...']	['great', 'phone', 'replace', 'lost', 'phone', 'only...']

```
In [4]: df = data[['Reviews', 'sentiment_encoded']]
df.head()
```

Out[4]:

	Reviews	sentiment_encoded
0	i feel so lucky to have found this used (phone...	2
1	nice phone, nice up grade from my pantach revu...	2
2	very pleased	2
3	it works good but it goes slow sometimes but i...	2
4	great phone to replace my lost phone. the only...	2

In [5]: `df.shape`

Out[5]: (349470, 2)

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 349470 entries, 0 to 349469
Data columns (total 2 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Reviews                349470 non-null object
1   sentiment_encoded      349470 non-null int32
dtypes: int32(1), object(1)
memory usage: 4.0+ MB
```

In [7]: `#Test and Train dataframes`

```
In [8]: %%time
import pandas as pd
from sklearn.model_selection import train_test_split

# Assuming 'data' is your DataFrame containing 'Reviews' and 'sentiment_encoded' columns
df = data[['Reviews', 'sentiment_encoded']]

# Split the data into training (60%) and temporary data (40%)
X_train_temp, X_temp, Y_train_temp, Y_temp = train_test_split(df['Reviews'], df['sentiment_encoded'],
                                                                test_size=0.4, random_state=42)

# Split the temporary data into testing (50%) and validation (50%)
X_test, X_validation, Y_test, Y_validation = train_test_split(X_temp, Y_temp, test_size=0.5,
                                                              random_state=42)

print("Train:", X_train_temp.shape, Y_train_temp.shape)
print("Test:", X_test.shape, Y_test.shape)
print("Validation:", X_validation.shape, Y_validation.shape)

Train: (209682,) (209682,)
Test: (69894,) (69894,)
Validation: (69894,) (69894,)
CPU times: total: 15.6 ms
Wall time: 69.9 ms
```

In [ ]:

In [9]: `#Building a model`In [10]: `from simpletransformers.classification import ClassificationModel`

```
# Create a TransformerModel
```

```
model = ClassificationModel('bert', 'bert-base-cased', num_labels=3, args={'reprocess_
```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-cased and are newly initialized: ['classifier.weight', 'classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
In [23]: train_df = pd.DataFrame({
          'text': X_train_temp[:20000].replace(r'\n', ' ', regex=True),
          'label': Y_train_temp[:20000]
        })

        test_df = pd.DataFrame({
          'text': X_test[:8000].replace(r'\n', ' ', regex=True),
          'label': Y_test[:8000]
        })

        validation_df = pd.DataFrame({
          'text': X_validation[:8000].replace(r'\n', ' ', regex=True),
          'label': Y_validation[:8000]
        })
```

```
In [12]: train_df.shape
```

```
Out[12]: (20000, 2)
```

```
In [13]: test_df.shape
```

```
Out[13]: (8000, 2)
```

```
In [14]: validation_df.shape
```

```
Out[14]: (8000, 2)
```

```
In [15]: model.train_model(train_df)
```

C:\Users\SACHIN\anaconda3\Lib\site-packages\simpletransformers\classification\classification\_model.py:612: UserWarning: Dataframe headers not specified. Falling back to using column 0 as text and column 1 as labels.

```
warnings.warn(
  0%|          | 0/20000 [00:00<?, ?it/s]
Epoch:  0%|          | 0/1 [00:00<?, ?it/s]
Running Epoch 0 of 1:  0%|          | 0/2500 [00:00<?, ?it/s]
(2500, 0.4227670201926492)
```

```
Out[15]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [17]: result_test, model_outputs_test, wrong_predictions_test = model.eval_model(test_df)
```

```
C:\Users\SACHIN\anaconda3\Lib\site-packages\simpletransformers\classification\classification_model.py:1454: UserWarning: Dataframe headers not specified. Falling back to using column 0 as text and column 1 as labels.
  warnings.warn(
0%|          | 0/8000 [00:00<?, ?it/s]
Running Evaluation: 0%|          | 0/1000 [00:00<?, ?it/s]
```

In [ ]:

In [24]: `result_validation, model_outputs_validation, wrong_predictions_validation = model.eval`

```
C:\Users\SACHIN\anaconda3\Lib\site-packages\simpletransformers\classification\classification_model.py:1454: UserWarning: Dataframe headers not specified. Falling back to using column 0 as text and column 1 as labels.
  warnings.warn(
0%|          | 0/8000 [00:00<?, ?it/s]
Running Evaluation: 0%|          | 0/1000 [00:00<?, ?it/s]
```

In [ ]: `#Model Evaluation test`

In [26]: `result_test`

Out[26]: `{'mcc': 0.7407787985957375, 'eval_loss': 0.35076412666449325}`

In [27]: `model_outputs_test`

Out[27]: `array([[ -3.09651875, -1.62132001, 4.50514841],
 [-1.52153051, -0.06079921, 1.81366718],
 [ 2.50336242, -0.79744315, -1.93051732],
 ...,
 [-2.14424634, -1.27033639, 3.46602249],
 [-3.15373158, -1.51070786, 4.44351864],
 [ 0.44362289, 0.6996659 , -0.63016003]])`

In [28]: `lst_test = []
for arr in model_outputs_test:
 lst_test.append(np.argmax(arr))`


In [29]: `true_test = test_df['label'].tolist()
predicted_test = lst_test`

In [30]: `import sklearn
mat_test = sklearn.metrics.confusion_matrix(true_test , predicted_test)
mat_test`

Out[30]: `array([[1749, 91, 127],
 [ 227, 106, 286],
 [ 125, 105, 5184]], dtype=int64)`

In [31]: `df_cm_test = pd.DataFrame(mat_test, range(3), range(3))

sns.heatmap(df_cm_test, annot=True)
plt.show()`



```
In [32]: sklearn.metrics.classification_report(y_true=true_test,y_pred=predicted_test,target_names=
```

```
Out[32]: '          precision    recall  f1-score   support\n\n      neutral      0.83246\n0.88917  0.85988      1967\nnegative   0.35099  0.17124  0.23018      619\npositive   0.92621  0.95752  0.94160     5414\naccuracy\n0.87987      8000\nmacro avg   0.70322  0.67264  0.67722     8000\nweighted avg   0.85865  0.87987  0.86646     8000'
```

```
In [33]: sklearn.metrics.accuracy_score(true_test,predicted_test)
```

```
Out[33]: 0.879875
```

```
In [34]: from sklearn.metrics import classification_report, accuracy_score, precision_score, re
```

```
# Calculate accuracy
accuracy = accuracy_score(true_test, predicted_test)

# Calculate precision, recall, and f1-score
precision = precision_score(true_test, predicted_test, average='weighted')
recall = recall_score(true_test, predicted_test, average='weighted')
f1 = f1_score(true_test, predicted_test, average='weighted')

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

```
Accuracy: 0.879875
Precision: 0.8586523308064559
Recall: 0.879875
F1 Score: 0.8664644261886704
```

```
In [35]: #Give your statement
```



```
In [85]: def sentiment(text):
    result = model.predict([text])
    pos = np.where(result[1][0] == np.amax(result[1][0]))[0][0]
    if pos == 0:
        print("This statement is Negative")
    elif pos == 2:
        print("This statement is Positive")
    elif pos == 1:
        print("This statement is Neutral")
```

```
In [86]: sentiment("a pretty capable, durable texting phone, unlocked as advertised, but the ve
0%|          | 0/1 [00:00<?, ?it/s]
0%|          | 0/1 [00:00<?, ?it/s]
This statement is Neutral
```

```
In [87]: sentiment("Do not buy it!. It feels like a pretended toy in your hand. If you want it
0%|          | 0/1 [00:00<?, ?it/s]
0%|          | 0/1 [00:00<?, ?it/s]
This statement is Negative
```

```
In [88]: sentiment("I went through lots of reviews for different phone before buying it. I four
0%|          | 0/1 [00:00<?, ?it/s]
0%|          | 0/1 [00:00<?, ?it/s]
This statement is Positive
```

```
In [41]: # validation
```

```
In [43]: lst_validation = []
    for arr in model_outputs_validation:
        lst_validation.append(np.argmax(arr))


    true_validation = validation_df['label'].tolist()
    predicted_validation = lst_validation

    import sklearn
    mat_validation = sklearn.metrics.confusion_matrix(true_validation , predicted_validation)
    mat_validation
```

```
Out[43]: array([[1792,   98,  148],
    [ 213,  129,  252],
    [ 140,  110, 5118]], dtype=int64)
```

```
In [72]: df_cm_validation = pd.DataFrame(mat_validation, range(3), range(3))

    sns.heatmap(df_cm_validation, annot=True)
    plt.show()
```



```
In [73]: sklearn.metrics.classification_report(y_true=true_validation,y_pred=predicted_validation)
```

```
Out[73]:
```

	precision	recall	f1-score	support	neutral	0.83543
0.87929	0.85680	2038	negative	0.38279	0.21717	0.27712
positive	0.92751	0.95343	0.94029	5368	accuracy	594
0.87987	8000	macro avg	0.71524	0.68330	0.69140	8000
g	0.86361	0.87987	0.86978	8000	nweighted avg	

```
In [74]: from sklearn.metrics import classification_report, accuracy_score, precision_score, recall_score
```

```
# Calculate accuracy
accuracy = accuracy_score(true_validation, predicted_validation)

# Calculate precision, recall, and f1-score
precision = precision_score(true_validation, predicted_validation, average='weighted')
recall = recall_score(true_validation, predicted_validation, average='weighted')
f1 = f1_score(true_test, predicted_validation, average='weighted')

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

```
Accuracy: 0.879875
Precision: 0.8636074021626677
Recall: 0.879875
F1 Score: 0.527322826895247
```

```
In [78]: import pandas as pd
```

```
# Assuming you have three DataFrames: train_df, validation_df, and test_df

# Export train_df to CSV
train_df.to_csv('train_data.csv', index=False) # Use index=False to exclude the index
```

```
# Export validation_df to CSV
validation_df.to_csv('validation_data.csv', index=False)

# Export test_df to CSV
test_df.to_csv('test_data.csv', index=False)
```


```
In [82]: import matplotlib.pyplot as plt
import numpy as np

# Data
models = ["Logistic Regression BOW BIGRAM", "Random Forest Regression TF-IDF",
          "Linear SVM Regression TF-IDF", "Neive Bayes Regression BOW BIGRAM", "BERT"]
accuracy = [0.8938, 0.9231, 0.8822, 0.8836, 0.8799]

# Define light colors
colors = ['#66c2a5', '#fc8d62', '#8da0cb', '#e78ac3', '#a6d854']

# Create a horizontal bar plot
plt.figure(figsize=(8, 4)) # Adjust the figure size as needed
plt.barh(models, accuracy, color=colors)
plt.xlim(0.75, 1) # Set the x-axis range
plt.xlabel('Accuracy')
plt.title('Model Accuracy')
plt.gca().invert_yaxis() # Invert the y-axis for better readability

# Show the plot
plt.show()
```



```
In [ ]:
```